HPC (V-1)

DOMS | Page No. | Date / /

Parallel Processing Computing

→ • HPC → high performance computing. (abstract level understanding)
• performing computational operations collaboratively → multiple computers.

→ Need → perform a high no. of operations per seconds (FLOPS)
→ complete a time-consuming operation in less time.
→ complete an operation under a tight dealine.
→ Handle huge amt of data.

→ Parallel computing → simultaneous use of multiple compute resources to solve a computational problem.

Appln → • Weather Forecasting • Scientific simulations • Big data analytics • Artifical Intelligence • Computational Biology • Financial Modeling.

Feature	★ Stored Program Architecture	★ General Purpose Cache Base Microprocessor Architect
Memory Access.	• all instructions & data fetched from main memory.	• fetched from cache or main memory
complexity	• Simpler Architecture	• More complex due to cache management
Performance	• Limited by main memory access speed	• improved due to faster cache access
versatality	• versatile, can execute diff programs easily	• optimized for specific tasks (not support diverse tasks)
scalability	• May face challenges in scaling for performance.	• scalable by increasing cache sizes or adding levels
examples →	• early mainframe computers	• modern personal computers, laptops.

DOMS | Page No.
Date / /

Motivating Parallelism –

- accelerating computing speeds
- multiplicity of datapaths
- increased access to storage elements.
- scalable performance & low cost

Implicit Parallelism –

- allows programmers to write their programming without any concern about exploitation of parallelism.
- compiler, runtime system & underlying hardware → imp role in exploiting parallelism implicitly.
- parallelism → transparent to programmer
  ↳ write standard sequential program without adapting any special parallel constructs
- underlying system → figure out parallelism from sequential code & help with diff techniques.
- present processor → use of multiple functional units & complete several instruction in similar cycle
- accurate method → commands are selected & execute provide remarkable diversity in architecture

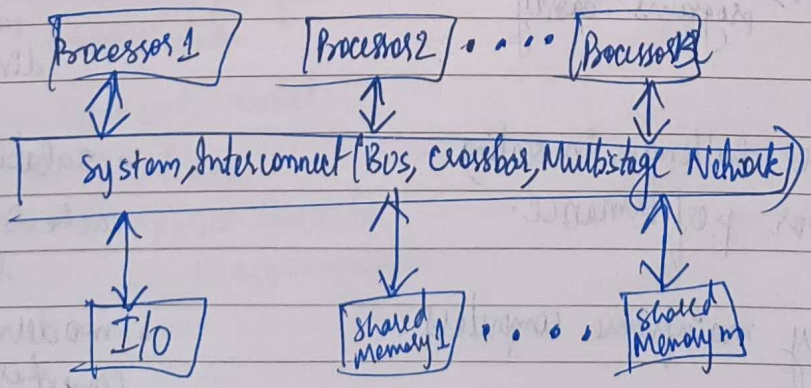

DOMS | Page No.
Date / /

✱ Dichotomy of Parallel Computing Platforms -

- openly parallel program → must agree concurrency & communication b/w simultaneous subtasks.
- previous → organize structure . latter → message model.
- logical & ~~physical~~ physical organization of parallel platform.
  - logical → (programmer's view of platform)
  - physical → (actual hardware side organization of platform)

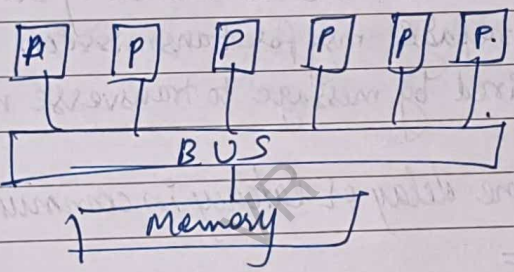

(Dichotomy shared memory model)

✱ Physical Organization of Parallel Platforms.

- Ideal Parallel Computer → PRAM (Parallel Random Access Machine)
  - Random Access Machine (RAM) → simple model of computation.
  - memory → unbounded sequence of registers.
  - each register → integer value.
  - control unit of RAM → program.
  - programer counter → determines statement to be executed next.
  - processes in PRAM → share same address space.

PRAM → 4 subclass → pattern of memory access →
- Exclusive-read, exclusive write (EREW) PRAM
- concurrent-read, exclusive write (CREW) PRAM
- Exclusive-read, concurrent-write (ERCW) PRAM
- concurrent-read, concurrent-write (CRCW) PRAM

DOMS | Page No.
Date / /

Communication Costs of Parallel Machines

→ processing elements leads to major overhead in parallel computing.

→

Cost of communication → depends on features
- Programming modes semantic
- Network Topology
- Data handling & routing
- Software protocols associate to program.

Message Passing Costs within Parallel Computers

→ Commⁿ time b/w nodes is characterized by sum of -
- Time to prepare msg for transmission
- Time required by message to transverse network to dest node

Parameters determine delay or latency in communication.

1) Startup time (t_s) =
- time consumed at sending & receiving nodes.

2) Per-hop time (t_h) =
- time taken by header of a message to travel b/w two directly connected nodes in network.
- also called as Node Latency
- t_h depends on delay in routing switch

3) Per word transfer time (t_w) -
- time required by each word to transverse link.
- $t_w = \frac{1}{r}$ → r words transferred through channel if bandwidth is r words/sec.

Parallel Machines (Two routing Techniques)

1)	Store & Forward Routing	Packet Routing	Cut Through Routing.
Concept →	• entire packet is received & stored before forwarding	• Forward based on packet destination address	• Forwarding begins upon receiving packet header
Focus →	• Reliability & error correction	• efficient data movement	• Low latency for small msgs.
Adv →	• Simpler Implementation • Error detection/correction at each hop.	• Flexible for various techniques	• efficient bandwidth usage
Disadv →	• high latency for small msg • less efficient	• requires storage for header info.	• Increased complexity delivery, potential buffer- • not suitable for large msg.
Latency →	• High latency ↳ storing entire packet	• Moderate latency ↳ routing decision.	• Lowest latency ↳ as begins upon header receipt.
performance –	Reliable but may suffer from latency.	• Blanced performe/scalability	• efficient performance with low latency

• $t_{comm} = t_s + (mt_w + t_h)$

$t_{comm} = t_s + mt_w$

• $t_{comm} = t_s + t_{w_1}m + t_{h1} + t_{w_2}(r+s) + (\frac{m}{s} - 1)t_{w_2}(r+s)$

• $t_{comm} = t_s + 1t_h + t_w m$

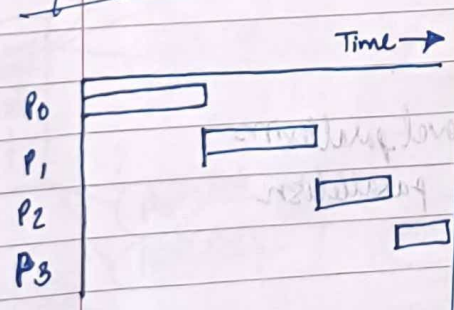

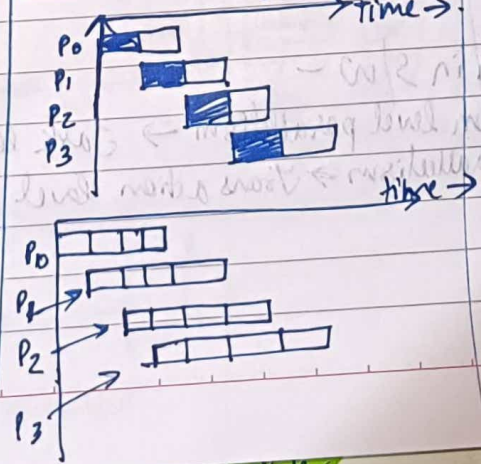

Cost Model for Communicating Messages

→ communicate message b/w two nodes l hops away using cut-through routing.

Total cost → $t_{comm} = t_s + lt_h + t_{wm}$.

To optimise cost
- communicate in bulk → small msg into single large msg
- minimize vol. of data → cost of per word transfer time can be reduced.
- Minimize distance of data transfer → min. no. of hops.

cost b/w two nodes → $t_{comm} = t_s + t_{wm}$

Levels of parallelism
- Instruction level parallelism
- thread level / task level parallelism
- Data level parallelism.

Parallism in Hardware (Uniprocessor):
- Pipelining, - Superscalar, VLIW etc.
- SIMD instructions, Vector processors, GPU.

Multiprocessors -
- symmetric shared memory multiprocessors
- Circulated " "
- Chip multiprocessor aka Multicores
- Multi computer a.k.a clusters.

Parallelism in S/W -
- Instruction level parallelism → Task level parallelism
- Data parallelism → Transaction level parallelism

(1-9, 1-22)

DOMS | Page No. | Date / /

Models – (SIMD, SIMT, SMIMD, SPMD)

SIMD	MIMD	SIMT	SPMD
• Single Instruction Multiple data.	• Multiple Instruction Multiple data.	• Single Instruction on Multiple Threads	• Single Program Multiple data.
• single instruction to multiple data element simultaneously	• Multiple inst^{n} to multiple data elements simultaneously	• singe instruction issued to multiple threads	• single program executed by multiple processors
• all data elements ↳ should have same datatype & operations	• data elements ↳ diff data types & operations	• all threads → diff data elements	• each processor ↳ operates on own data partition
• easier to program for specific tasks.	• more complex ↳ data dependencies & race conditions	• similar to SIMD allows for thread divergence	• similar to MIMD, focusses on data parallelism with a single program.
• minimal comm^{n} b/w processing units/ elements.	• comm^{n} may occur for coordination or data sharing	• comm^{n} within thread groups for synchronization	• comm^{n} may occur.
• High data level parallelism	variable data level parallelism	Moderate data level parallelism within thread group	Moderate data level parallelism within same program execution
• eg → vector processors, GPU.	eg → distributed computing, clusters	eg → CPU architecture, CUDA.	eg → MPI, OpenMP

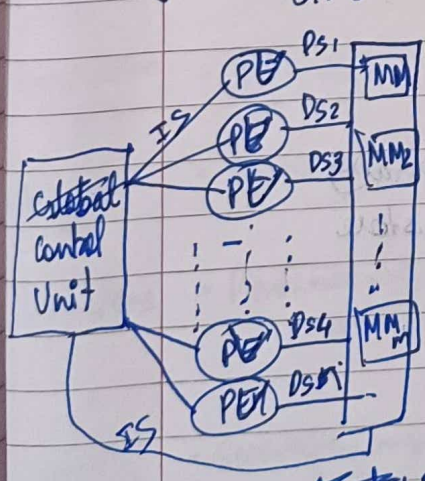

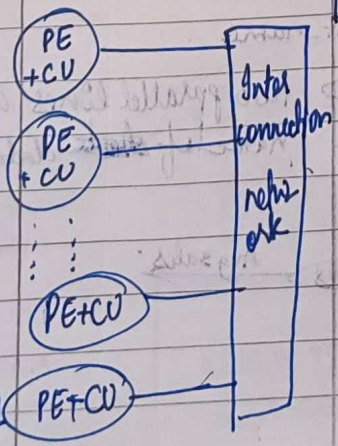

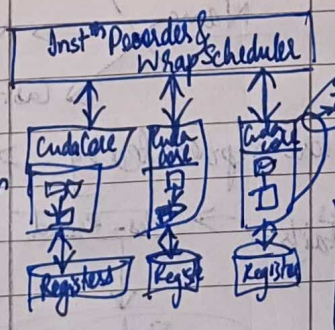

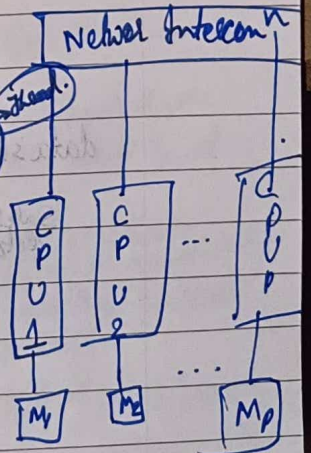

PE → Processing ~~Element~~ Unit
MM → memory module
DS → Data Stream.

CU → control Unit.

DOMS | Page No.
Date / /

Data Flow Models –

- defines fun" of internal process system with aid of Data Flow Diagram
- depicts function derivation of data values without indicating how they derived.

Data Flow Diagrams –

- Functional modelling → Hierachy of DFD.
- graphical representation of system → shows input to system, processing upon inputs, the outputs of system as well as internal data stretches.
- illustrate series of transformation & computations performed.

4 parts of DFD → 1) Processes → • transform data values
• visualised as high level process → system
2) Data Flows → Flow of data b/w processes
↳ the details process or data of process store
3) Actors → active objects interact with system
4) Data Stores → passive object act as repositories of data

other parts → constraints & ~~data~~ control flows

data flow represented by → directed arc or an arrow, labelled with name of data item.

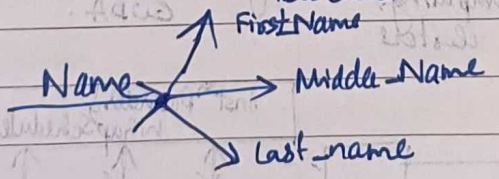

data store represented by → two parallel lines containing name of ~~data~~ data_store

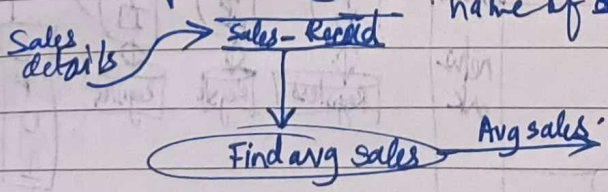

DOMS | Page No.
Date / /

* Demand Driven Computation –

- general framework for deriving demand driven algorithm for interprocedural data flow analysis.
- goal → reduce time &/or space overhead of conventional exhaustive analysis by avoiding collection of info
- Reduction Machines/Lazy Computers.
- approach matches with functional programming language.
- whether minimizing time or space is of primary concern. result caching may be incorporated in derived algo.
- If problem's flow function → distributive → derived algo provides as precise info as corresponding exhaustive analysis
- executed when results are required.
- NO use of shared memory.
- High degree of parallelism.

Appln
- Database servers, web servers, embedded network, digital signal processing, multimedia application, scientific application & thread level parallelism

Feature	N-Wide Superscalar Archit	Multicore Architecture	Multithreaded Architecture
Processing units	• multiple execution units within a core	• Multiple independent cores	• single core with thread context switching
Parallelism Type	• Instruction level	• Task level (program)	• Thread level
Benefits	• Improved performance	• significant performance gains for parallelized workloads.	• Improved performance for thread level.
Limitations	• Diminishing returns with more units	• Increase complexity require multicore aware S/w.	• Limited parallelism compared to multicore, overhead of context switching
Typical Use cases	• Scientific Computing, Image/Video Processing	• General purpose computing, multitasking, server workloads	Multitasking, web browsing I/O bound tasks

Clip Board

DOMS | Page No.
Date / /

* Pipelining -

- Breakes down instruction execution into stages.
- Allows overlapping of instruction execution.
- Increases throughput by processing multiple instructions simultaneously
- Reduces overall execution time.
- Instructions progress through stages sequentially.

* Superscalar Execution -

- Utilizes multiple execution units with single processor core.
- executes mutiple instructions simultaneously.
- can exploit instruction level parallelism with a program.
- Dispatches instruction-level parallelism to mutiple execution units.
- acheives greater performance gain compared to pipeline alone

* VLIW (Very long Instruction Word) processor.

- executes mutiple instructions in parallel using single VLIW.

Process it follows → 1) Instruction Bundling 2) Static Scheduling
3) Parallel Execution 4) No dependency checking
5) effiency & performance 6) compiler dependency.

Adv → • reduced hardware complexity • high performance for parallel appln
Disadv → • relies on complex compiler for parallel instruction • less flexible code
Appl^n → • Digital signal Processing • Scientific computing • Embedded systems

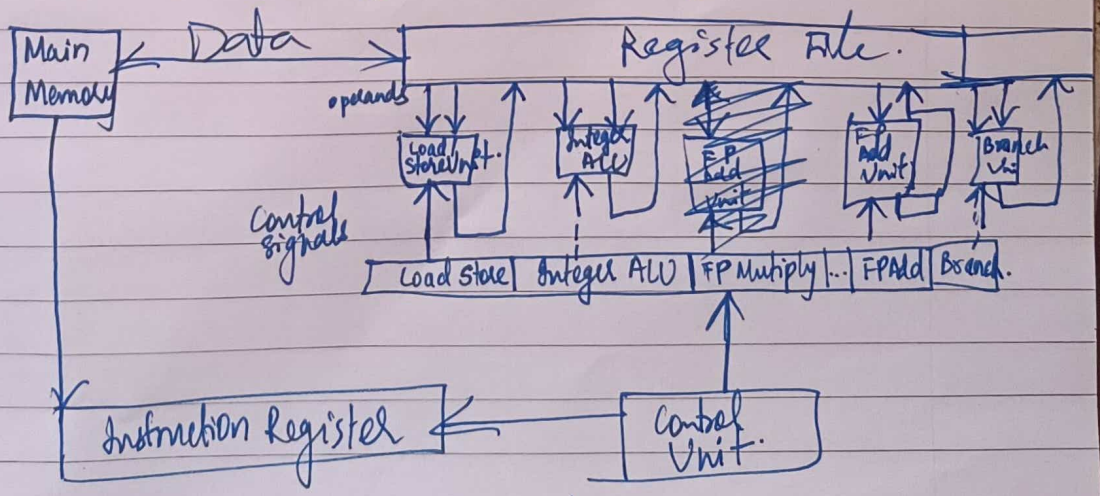

VLIW Architecture

DOMS | Page No.
Date / /

NUMA Multiprocessor

- Non uniform (varies by location)
- more complex, mutiple memory ctrls
- scalability is more
- performance higher (for workload with locality)
- require msg passing b/w processors
- eg → HPC, large DBs

UMA multiprocessor

- Unifor (equal access time)
- simples, simple mem ctrls
- limited scalability
- lower performance (potential bottleneck)
- Interprocess communication simpler due to shared memory
- eg → General purpose, time sharing